



PointNet for pflow

Joshua Himmens,
Dr. Luca Clissa,
Dr. Maximilian Swiatlowski

Outline

- The p-flow problem
 - Current approaches
- Using PointNet for p-flow
 - Why PointNet
- Processing tracks and cells
- Loss, Metrics, and baseline calculations
- Results
 - Several datasets
 - Modal collapse
- Ongoing work
 - Transfer learning
 - Multi-class labeling

Problem

Existing ATLAS particle flow (p-flow) algorithm

- 👍 calo & track information
- 👍 good energy and angular resolution
- 👍 very helpful with pile-up

Main limitations:

- 👎 associate track to topo-clusters, not cells directly
- 👎 hand-tuned "ordering" algorithm for cell energy subtraction
- 👎 fails when tracks are too dense

Goals



Test ML approaches to target cell-to-track association and replace hand-tuned ordering

The main idea is to start from single-particle MC data building upon [ATL-PHYS-PUB-2022-040](#)

Specific focus is be on:

- improve association performance
- extend to jet environments
- optimise neighborhood size

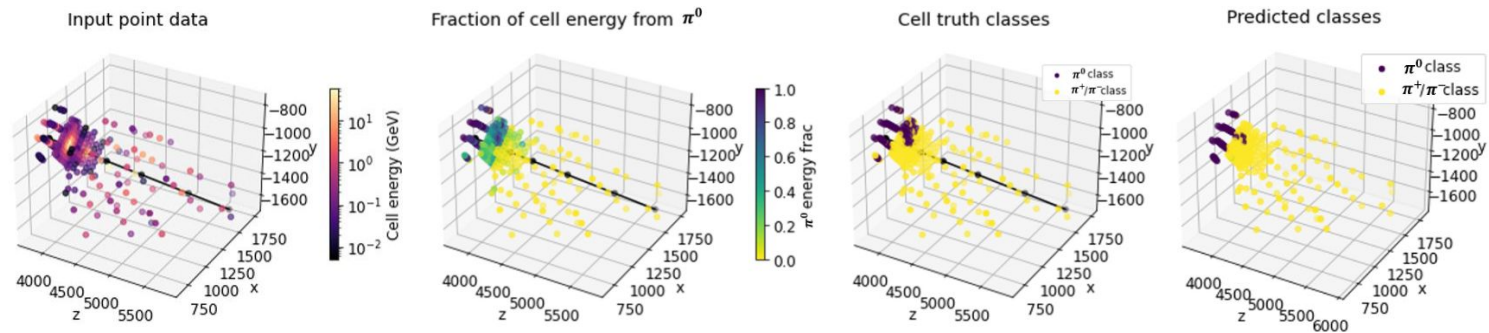
JIRA tracking: <https://its.cern.ch/jira/browse/ATLJETMET-1699>

Current strategy

Represent data as point clouds

- sparse, efficient representation
- enable joint use of calorimeter and tracker information

Adopt segmentation models, i.e. assign track labels to each point (track & calo)

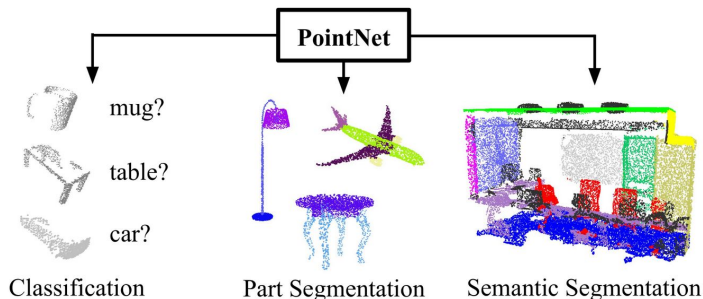


PointNet model [12]

The idea is to adopt NN architecture suitable for point clouds → **PointNet**

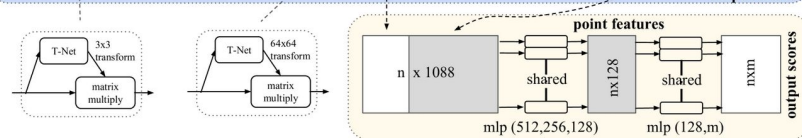
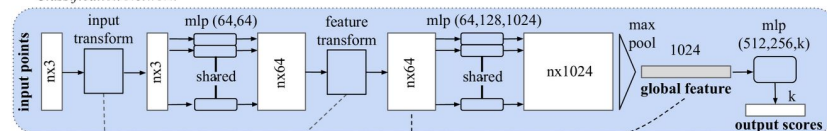
Several learning tasks:

- classification
- part segmentation
- semantic segmentation



- 👍 permutation invariant
- 👍 transformation equivariance
- 👍 both shape classification & segmentation
- 👍 robust to data corruption → critical points
- 👎 no local context → global feature learning
- 👎 generalization to unseen scenes → global features depend on absolute coordinates
- 👎 no rotation/shape equivariance

Classification Network



Segmentation Network

Datasets

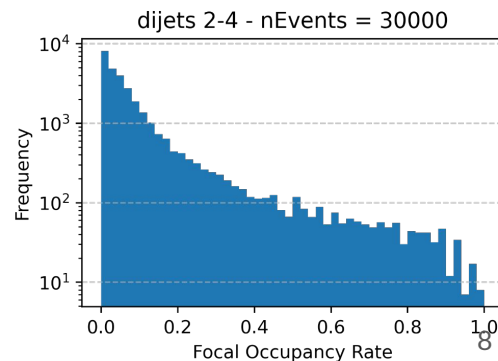
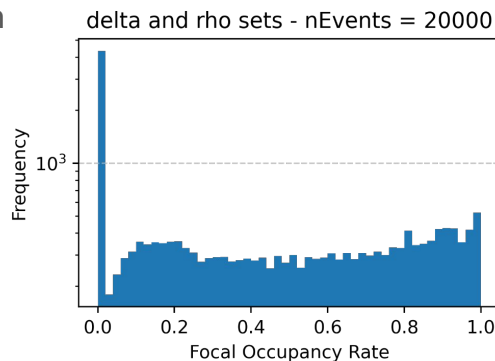
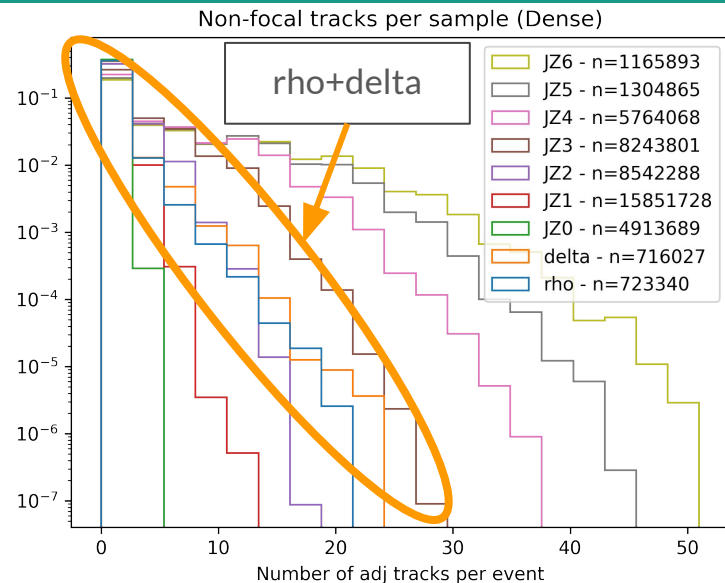
Heterogeneous Data

We analysed **several MC simulations** ([ML Tree](#)):

- **rho, delta decays**
 - typically 1 track per event
 - many focal calo hits per sample (~10-200)
- **dijets decays (split into JZ by pt)**
 - 1 to 32 tracks per event
 - few focal calo hits per sample (<50)

For each track, we create a *sample* with all hits in a ΔR of size 0.2

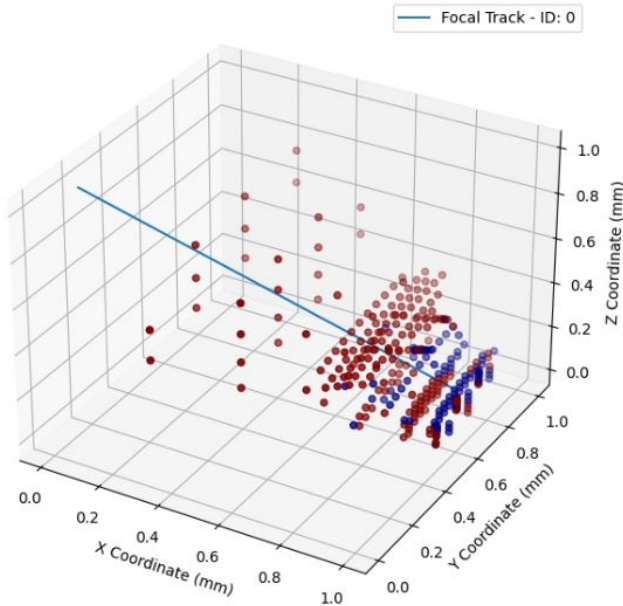
- the samples constitute individual point clouds fed into PointNet
- **focus VS “unfocus” hits**



Disparate Training Data - Example

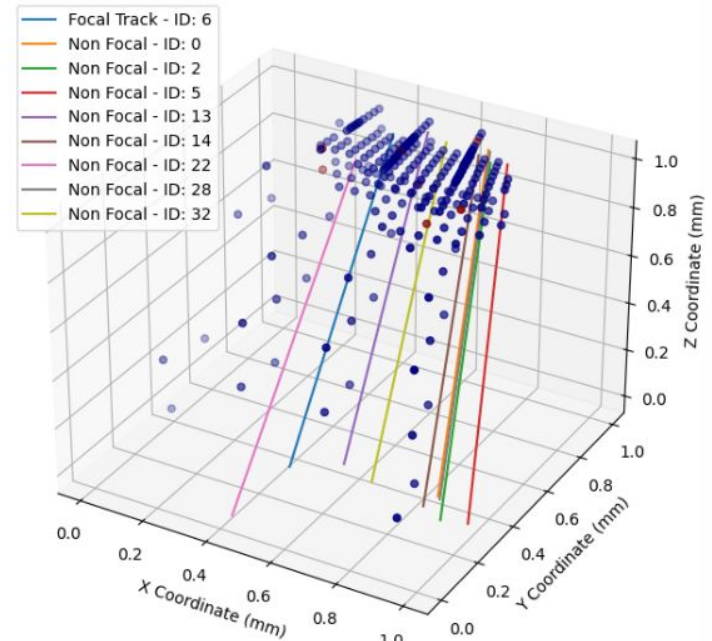
rho sample

Truth Values - 232 activations



dijet (JZ4) sample

Truth Values - 7 activations



Pre-processing



The majority of the complexity is the preprocessing to make the root data suitable for a PointNet.

- Flattening ROOT ntuples and cutting for ΔR
- Split events trackwise \rightarrow samples
- Remove erroneous tracks caused by simulation quirks
- Embed data into a matrix of points

All this is available as part of [JetPointNet on gitlab](#) under nightly.

.

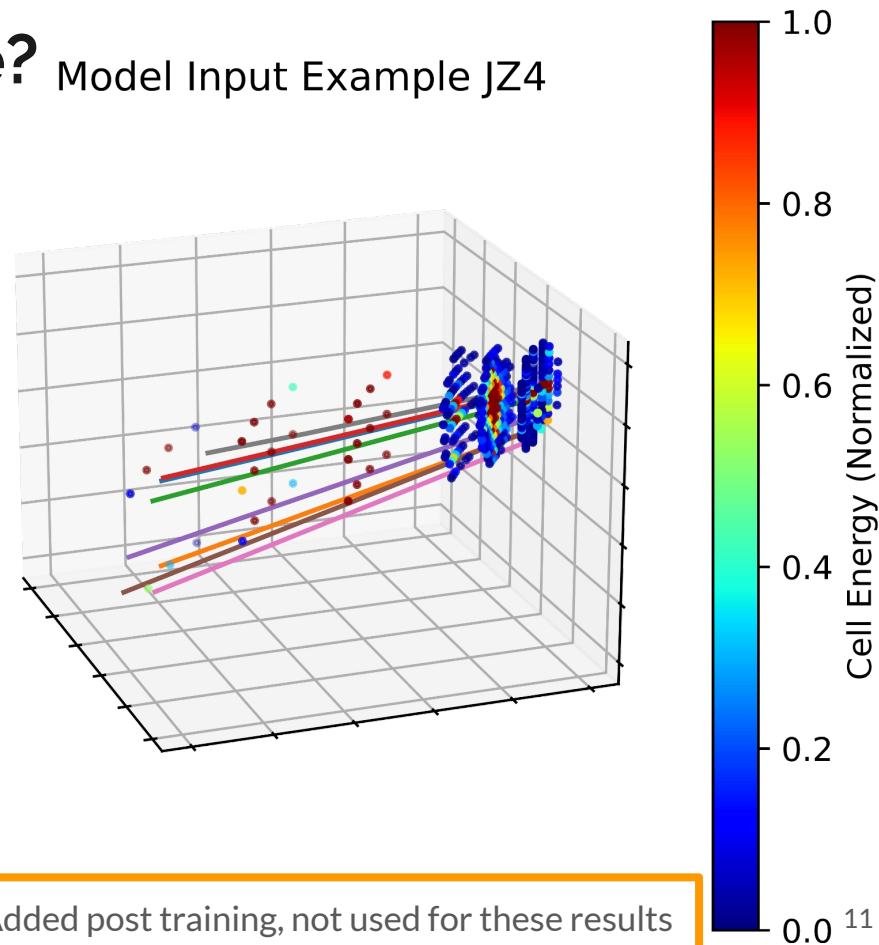
What does the model see?

Model Input Example JZ4

The cell hits and extrapolated track hit point clouds are overlaid.

Tracks are associated to a truth particle.*

Cells are attributed to a single truth particle even if they are a member of multiple clusters or have multiple contributing particles.



*There are cases where tracks have interactions in the tracker and lose their label, these events are cut from the training set.

*Added post training, not used for these results

Data Processing Challenges

Not all fields are relevant to all points.

Some points are semantically related, like track points, which is difficult to represent in a point cloud.

A static length input means that most events will be significantly padded.

*Added post training, not used for these results

Data encoded into each point
Data encoded into each point

Point type (e.g. cell hit, track, padding)

ΔR

Track Identifier (for tracks)

Normalized coordinates

Track χ^2/dof (for tracks)

Cell E (for cells)

Track pt (for tracks)

Cell Sigma (for cells)

Loss, Metrics and baseline

How do we measure performance?



Given the complexity of the problem, we must carefully what constitutes good performance and how to measure it.

- what do we care about?
 - mostly cell hits
 - more energetic hits
 - masked loss/metrics and energy weights
- how to measure it?
 - several alternative options, with pros and cons
 - no perfect metric, consider multiple together
 - data challenges: class imbalance (e.g., rho dataset: 66% class 0)
 - use class-weighted loss/resampling
 - accuracy alone is misleading

Loss



We adopt a standard **energy-weighted binary cross-entropy**:

$$\text{WBCE} = -\alpha * y \log(y^*) - (1-y) \log(1 - y^*)$$

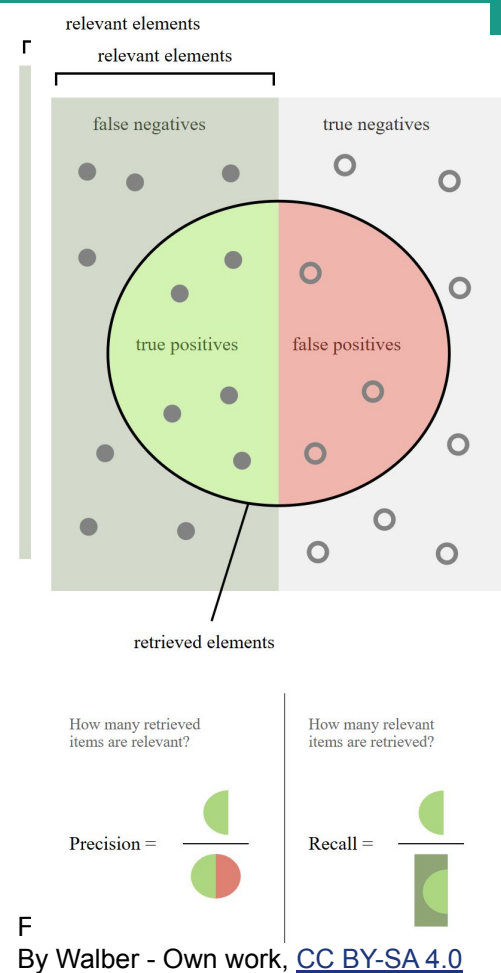
This should ensure that most energetic hits have larger impact on the loss, thus forcing the training to pay more attention to those hits.

Several weighting schemes have been attempted: none, quadratic, absolute value, Cox- transform....

Metrics

For a thorough assessment, we jointly look at several metrics:

- **Accuracy:**
 $(TP+TN) / (TP+TN+FP+FN)$
- **Precision** (purity):
 $TP / (TP + FP)$
- **Recall** (signal efficiency):
 $TP / (TP + FN)$
- **F1-Score:**
 $2 * P * R / (P + R)$



Baseline



How do we define “good” performance?

Baselines:

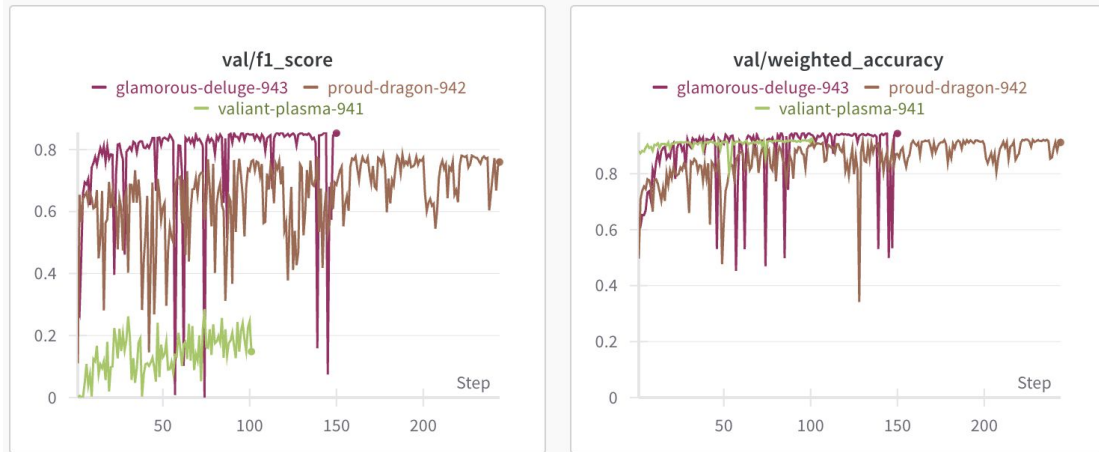
- “dumb” model: always predict majority class
- random model: random prediction (50/50 chance)
- “smart” random model: random prediction with probability linked to class proportions
 - *rho + delta*: ~64% of unfocus points (class 0)
 - JZ2 + JZ3 + JZ4: ~94% of unfocus points (class 0)

Model training

Training Infrastructure

We train using NVIDIA RTX A6000's and manage training with Weights and Biases.

To avoid long epochs, we randomly sample ~204,800 events (~14 minutes of compute). Model convergence requires 100 epochs (~23h). Our training data contains 246M trainable samples (1.1TB)



Example Weights and Biases Plots

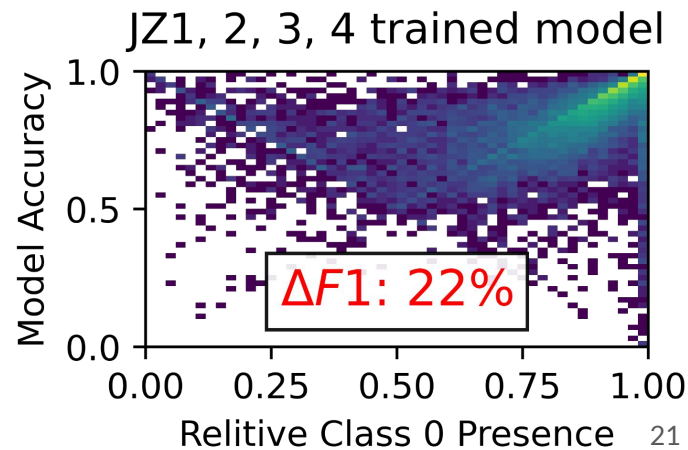
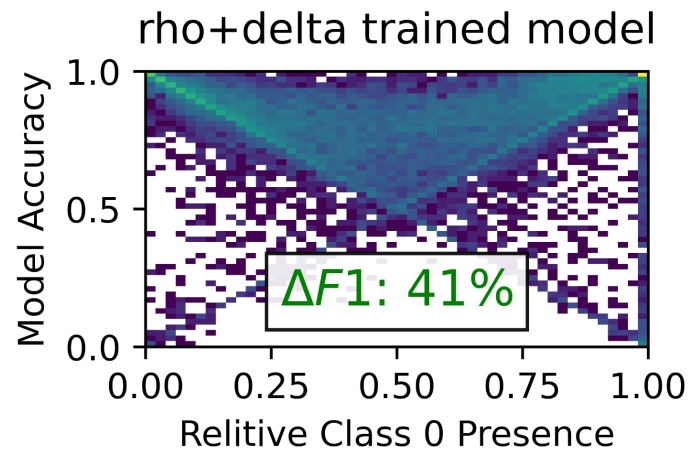
Results

Results

Training on the rho+delta sets is **successful** with an F1 score of ~0.8 and a significantly improved **energy weighted accuracy** >0.9.

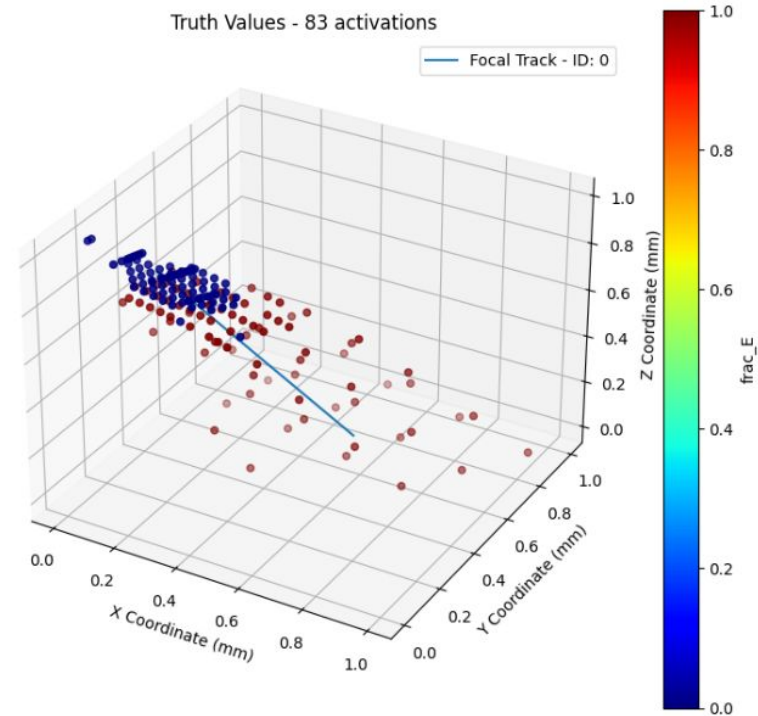
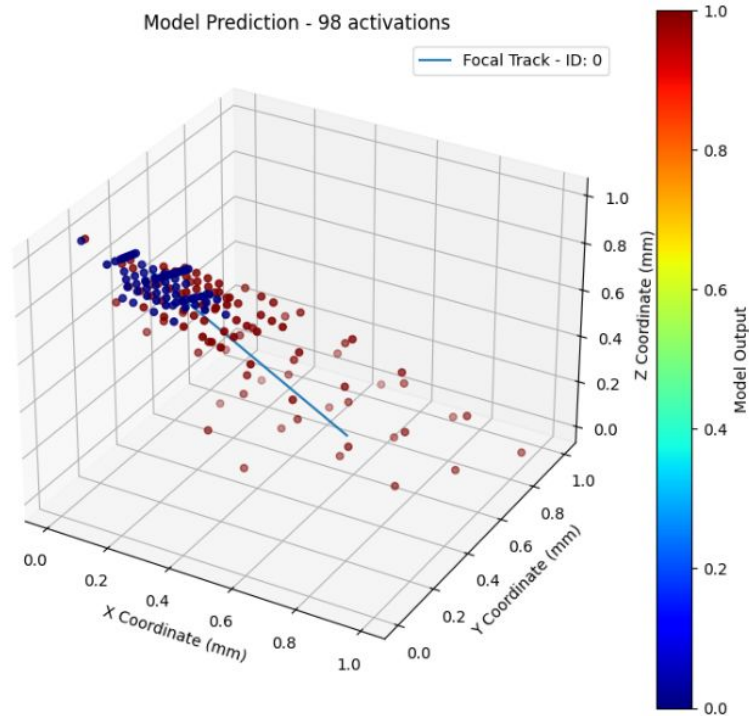
Dijet training is promising, however the model is **unable to overcome modal collapse**, where the entire event is attributed to non-focal tracks.

Dataset	Model	Accuracy	F1-score
<i>rho + delta</i>	<i>PointNet</i>	0.83	0.78
	<i>smart_random</i>	0.63	0.37
<i>JZ2 + JZ3 + JZ4</i>	<i>PointNet</i>	0.94	0.28
	<i>smart_random</i>	0.94	0.06



Example of a successful segmentation - rho

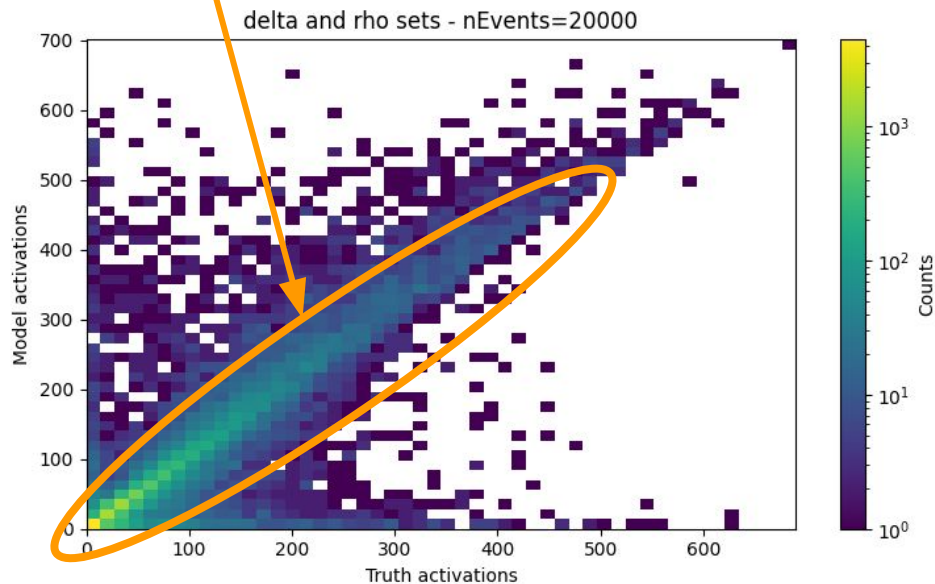
Model activated energy = 1179 GeV \approx Ideal Activation = 1159 GeV



Modal Collapse

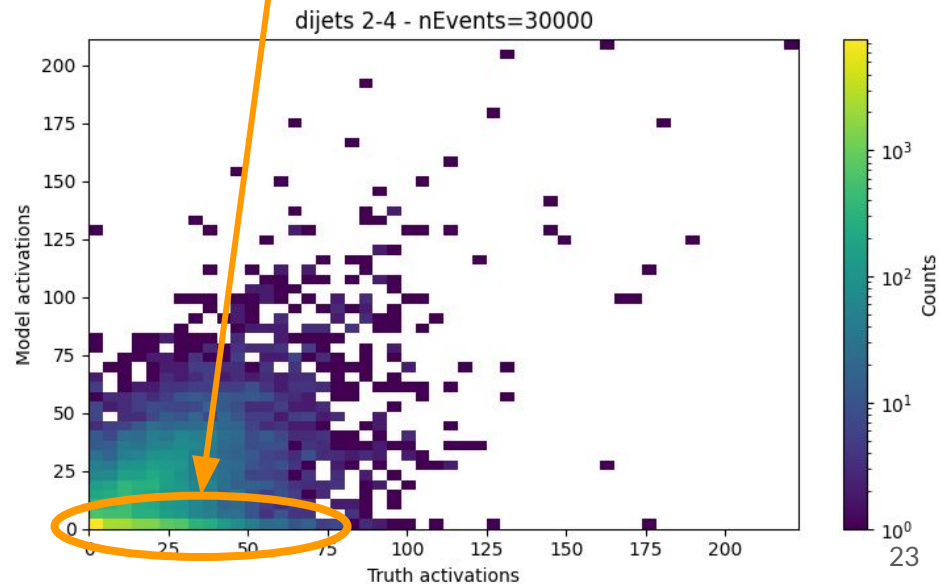
Models tend to become event-wise rather than cell wise for more complex sets.

Healthy Model



Truth v Model Activation

Large event-negative regime



Ongoing Work

Transfer Learning - Architectural

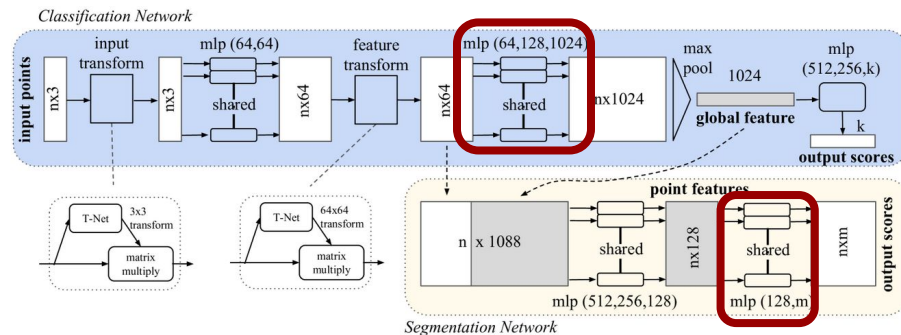
Attempting to transfer the positive results of the rho set to the dijet sets by freezing parameters to freeze model knowledge.



Classification Net: ~4M Params -> (200-500)k Params

Segmentation Net: ~1M Params -> (25-100)k Params

More hyper-parameter sweeping is needed.



Transfer Learning - Replay

Attempting to transfer the positive results of the rho set to the dijet sets by tuning with both the source and target sets.

Model appears to be a linear combination of a directly trained rho+delta and dijet model. More hyper parameter sweeping needed.

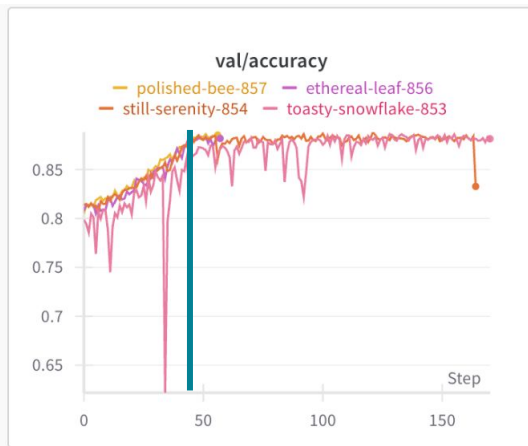
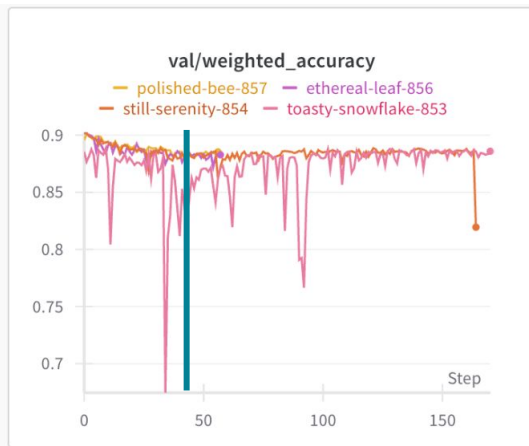
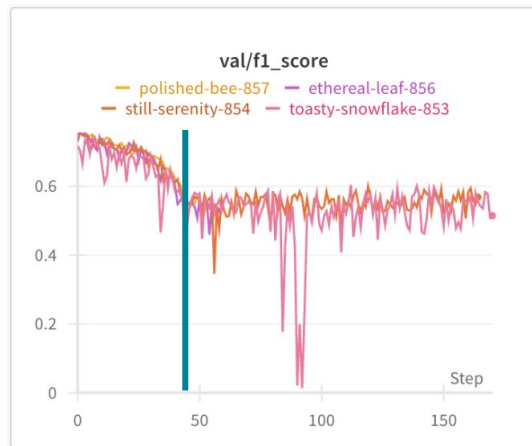
Strong performance in the rho set

Train to convergence
with rho dataset

Weaker performance in both sets (so far!)

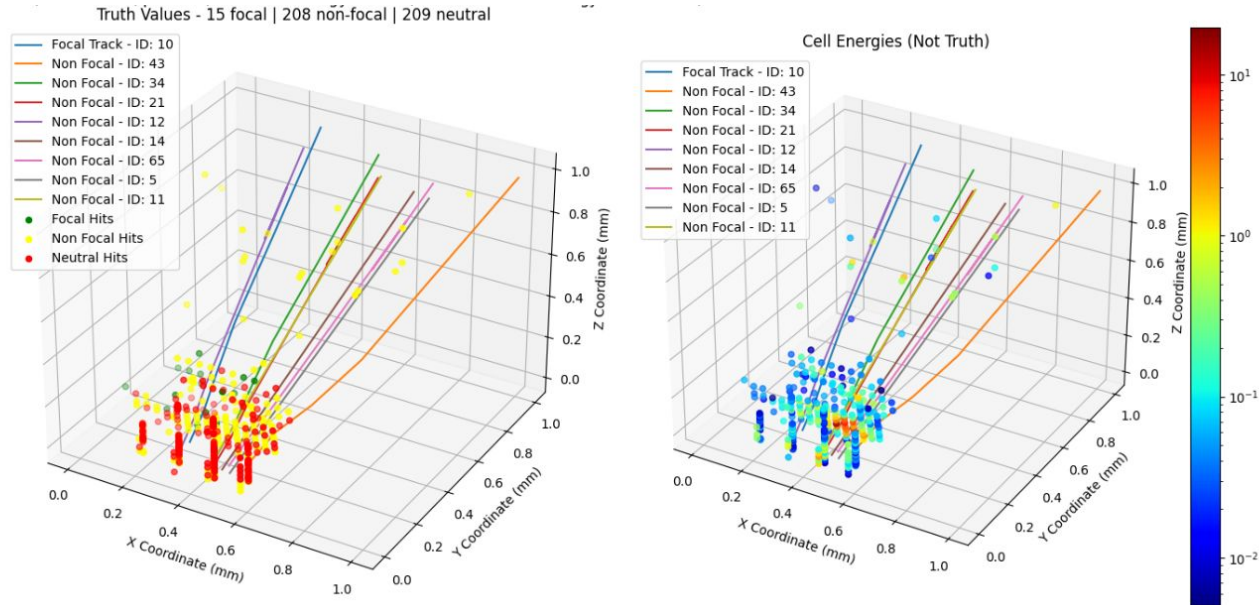
Train on rho set with
increasing presence of the dijet data

End of data injection phase



Multi-Class Segmentation to Reduce Confusion

Instead of binary class labels, use categorical information for focal, non focal, and neutral. We are hoping to **differentiate the negative class to reduce modal collapse**.

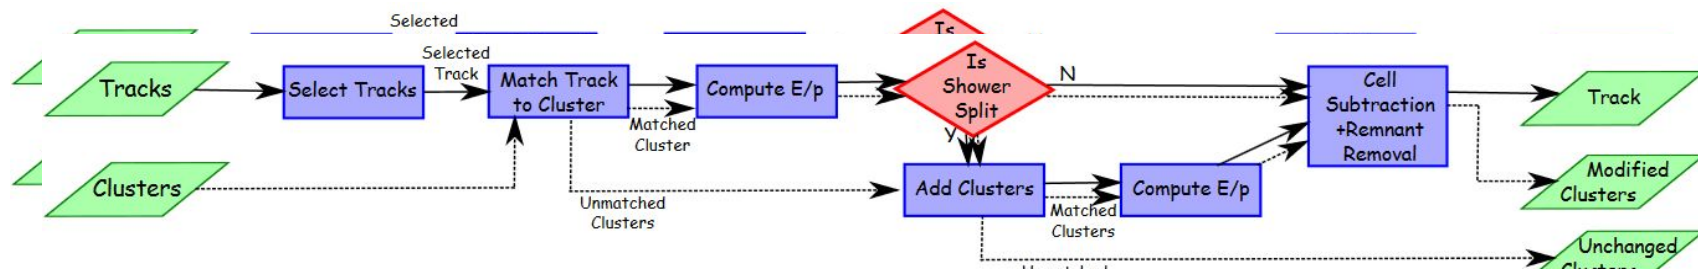


Backup

ATLAS p-flow algorithm [[ATL-PERF-2015-09](#)]

For track in descending pT:

- associate closest topo-cluster based on angular distance $\Delta R'$
- compute expected energy deposit based on the topo-cluster position and track momentum
- if expected and measured energies differ significantly, associate more topo-clusters
- subtract the expected energy by calo cells
- if remaining energy lies within expected fluctuations, remove the remnants



Baseline

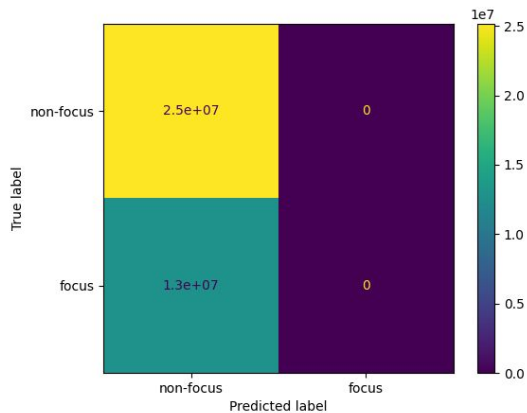
Whole datasets results considering all hits except padding

Dataset	Model	Accuracy	Precision	Recall	F1-score
<i>rho + delta</i>	<i>dumb</i>	0.6488	0.0000	0.0000	0.0000
	<i>random</i>	0.5001	0.3512	0.4999	0.4126
	<i>smart_random</i>	0.5443	0.3512	0.3511	0.3512
<i>JZ2 + JZ3 + JZ4</i>	<i>dumb</i>	0.9414	0.0000	0.0000	0.0000
	<i>random</i>	0.5000	0.0586	0.5000	0.1049
	<i>smart_random</i>	0.8897	0.0586	0.0586	0.0586

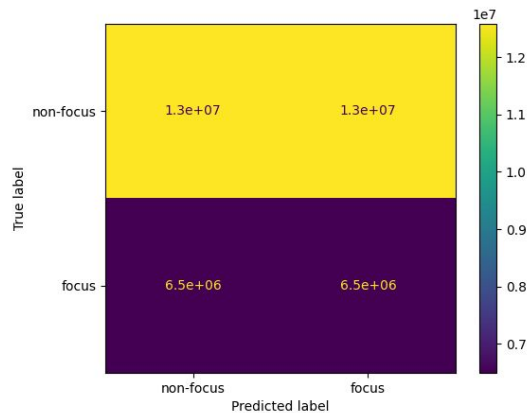
Confusion matrix



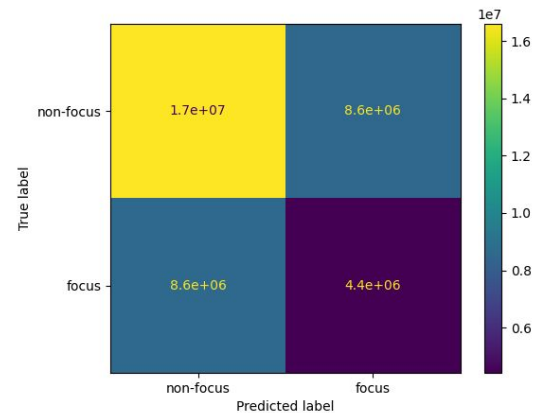
dumb model



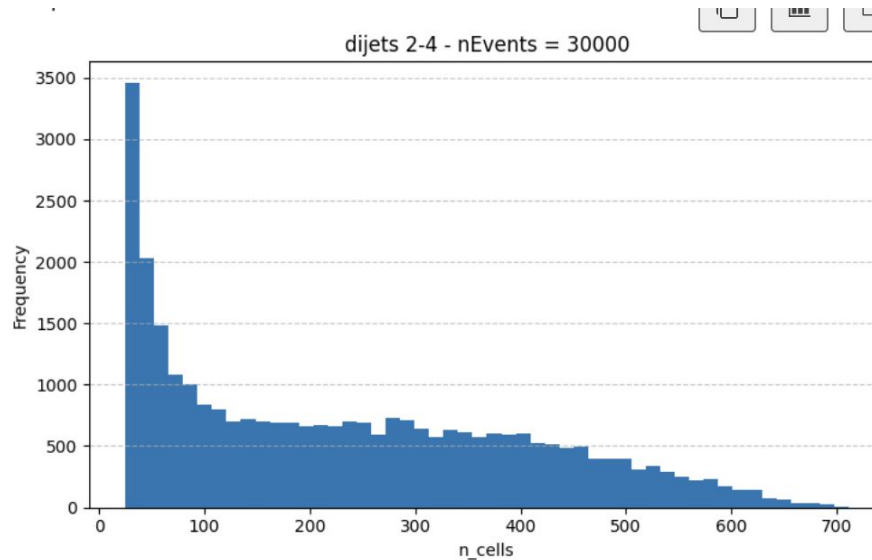
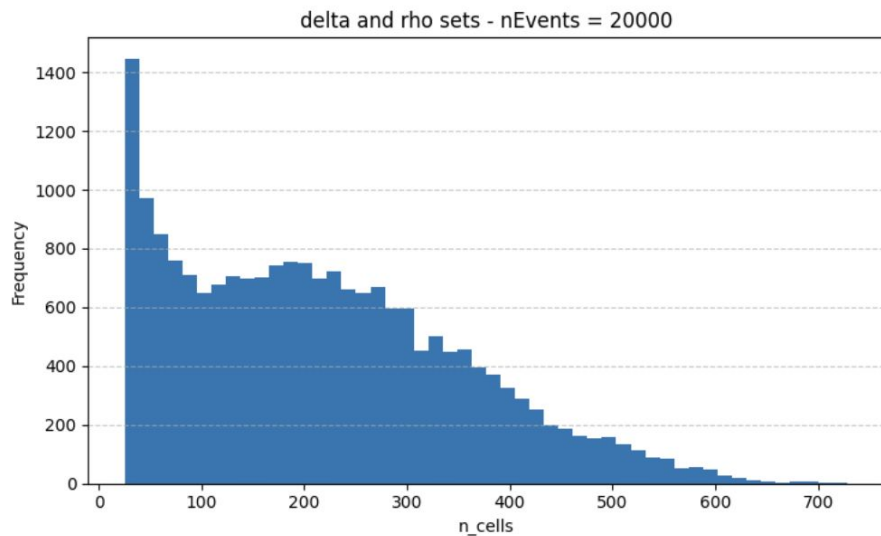
random model



smart random model



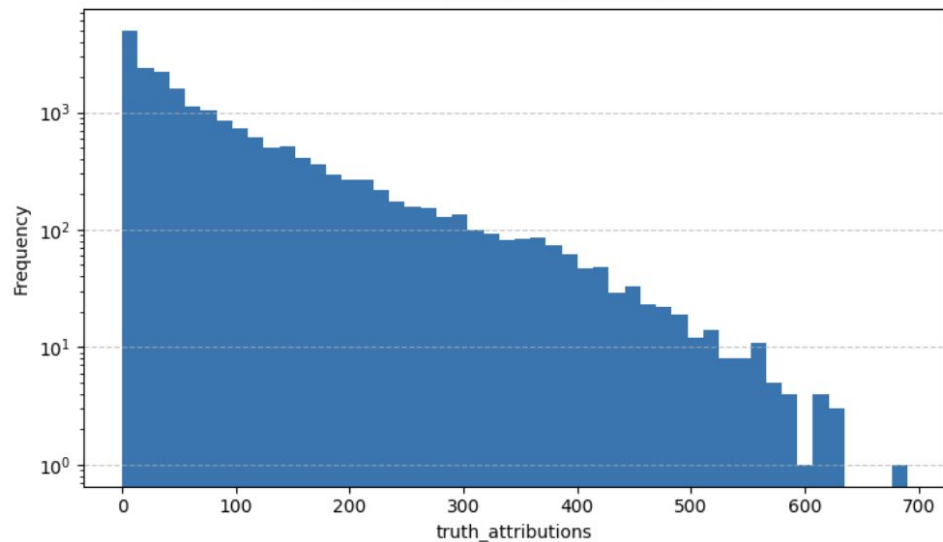
n_cells



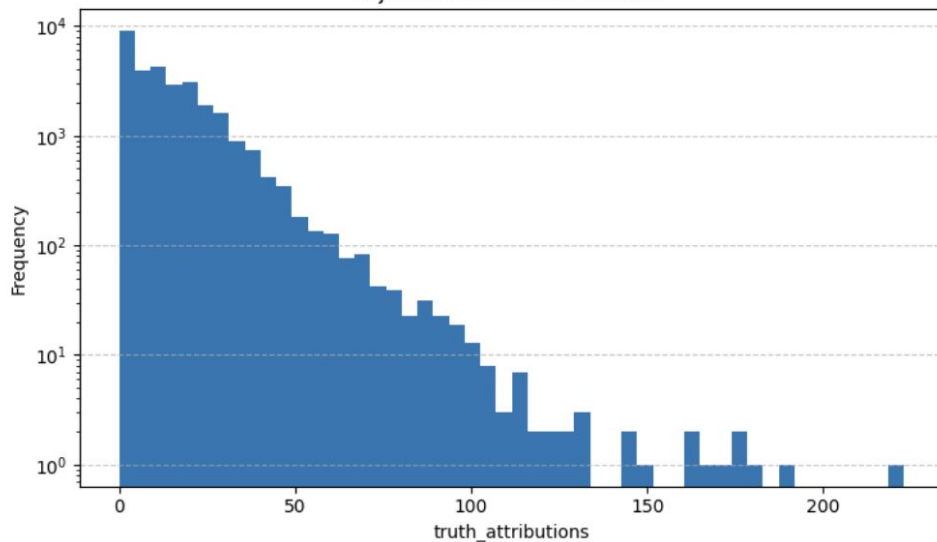
truth_attributions



delta and rho sets - nEvents = 20000



dijets 2-4 - nEvents = 30000





Dataset	Model	Accuracy	Statistical Certainty (F1)
<i>Simplified Data</i>	<i>PointNet</i>	0.83	0.78
	<i>Random Estimator</i>	0.63	0.37
<i>Realistic Data</i>	<i>PointNet</i>	0.94	0.28
	<i>Random Estimator</i>	0.94	0.06